# COMP20003 Algorithms and Data Structures Dictionaries and Data Structures Nir Lipovetzky Department of Computing and Information Systems University of Melbourne Semester 2

# So far... We have: Looked at algorithms, fast and slow. Estimated computation time by counting operations. Formalized a system for classifying algorithm efficiency.

## **Outline of the first few lectures**



- Algorithms: general
- This subject: details
- Algorithm efficiency: intuitive
- Computational complexity
- Data structures
  - Basic data structures
  - Algorithms on basic data structures
  - Complexity analysis of algorithms on basic ds's

20003 Algorithms and Data Structures

### **Textbook**



• Skiena: Chapter 3, Data Structures

COMP 20003 Algorithms and Data Structures

### This section



- A lightning tour of fundamental data structures used for search:
  - Arravs
  - Linked Lists
  - Trees

COMP 20003 Algorithms and Data Structures

# Abstract Data Types *vs.* Data Structures



- Abstract data type: what it does
  - Stack, queue.
  - · Dictionary: look up by key.
  - Does not specify an implementation
- Concrete data structure:
  - Array, linked list, tree
  - Implements an abstract data type.

1-6

### **Data structures**



- Organizing data is important
- It is helpful to organize with the task in mind
- For searching, e.g.:
  - Some high-level languages have inbuilt "dictionaries", or associative arrays (Python, awk).
  - In lower-level languages, dictionaries are implemented directly using a fundamental data structure.

03 Algorithms and Data Structures

### Searching



- Search Question:
  - Given a search key,
  - Find the record(s) that correspond to this key.
  - Typically we describe record simplistically with fields key and info (or just key).
- Examples:
  - Students and seat numbers, Telephone books

How you organize data is important for search!

### **Dictionaries**



Python built-in dictionary structure.

```
>>> tel = {'jack': 4098, 'sape': 4139}
>>> tel['guido'] = 4127
>>> tel
{'sape': 4139, 'guido': 4127, 'jack': 4098}
>>> tel['jack']
```

• The dictionary is implemented using one of the underlying data structures.

 sape	_		jack	guido		
4139			4098	4127		

### **Outline of the first few lectures**



- · Algorithms: general
- This subject: details
- Algorithm efficiency: intuitive
- Computational complexity
- Data structures
- Basic data structures
  - Algorithms on basic data structures
  - · Complexity analysis of algorithms on basic ds's

COMP 20003 Algorithms and Data Structures

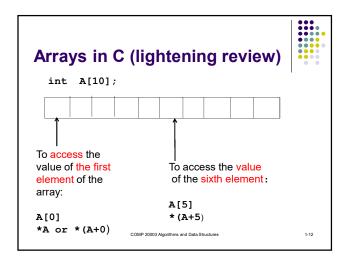
1-10

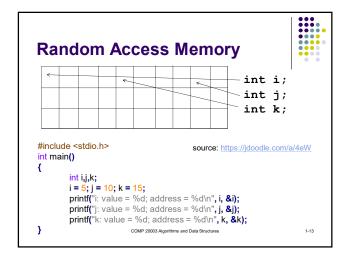
### **Array**

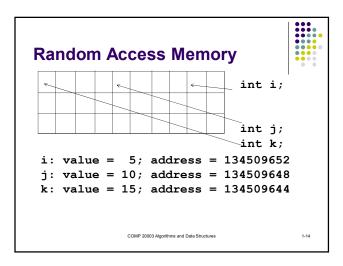


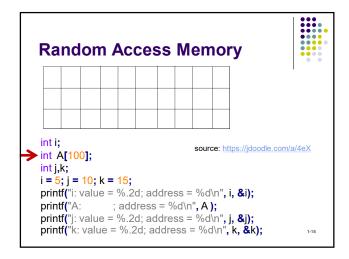
- Def: given an index (location), we can retrieve any item in unit time.
- If items are in arbitrary order, finding a key in array of size n requires O(n) time.

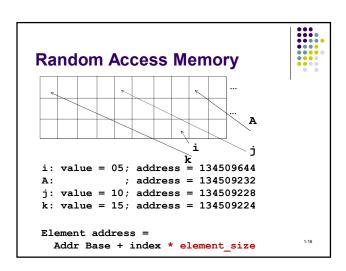
Mimics the structure of random access memory (RAM), where the index is the memory address

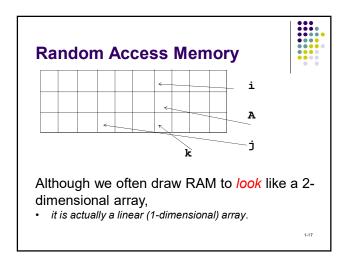


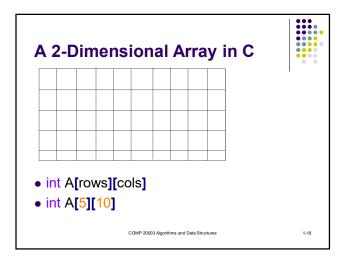


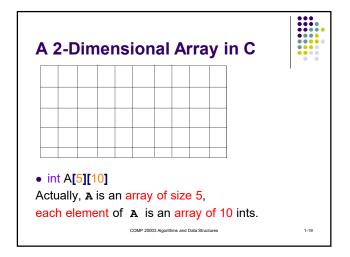


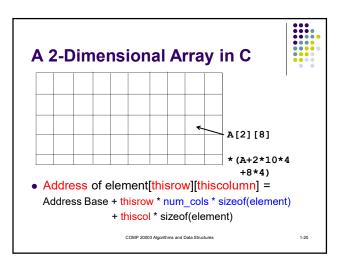


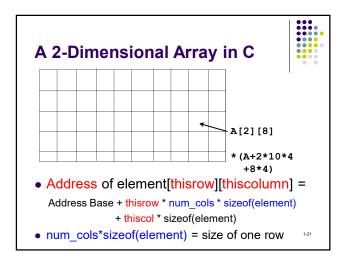












### **Homework**

- What is the difference between:
  - int a[10][20];
  - int \*b[10];
- ? 2d in array vs 1d pointer to int array
- Hint 1: Think memory allocation.
- Hint 2: See K&R section 5.9.
- Hint 3: How could you test your answer?

### Back to arrays as dictionaries



- To sort or not to sort?
- How to determine which is best?

Use big-O analysis!

1-23

### Sorting: fine points



- Sorting assumes the keys are "sortable",
   i.e. comparable. E.g.:
  - · Comparable: categories,
  - Not Comparable: colours, unless you associate an identifier.
- The CS definition of sorting means to put things into a *well-defined* order.
  - Other ways of sorting without comparing keys: counting sort. (next lectures)

### Sorting: fine points



- In CS applications of sorting, there is a key, and associated information.
- We sort by *key*, and the *information* comes along for the ride.
  - Student database, sorted on student ID.
  - Information: name, address, degree, etc.
- In our examples, we often do not show the *information* explicitly.

1-25

## **Unsorted arrays**



- Just put the item at the end of the array:
  - Insertion is in O(1)



 How many comparisons you need to insert? NONE







### 

### **Unsorted arrays**



- Just put the item in at the end of the array:
  - Insertion is in O(1)

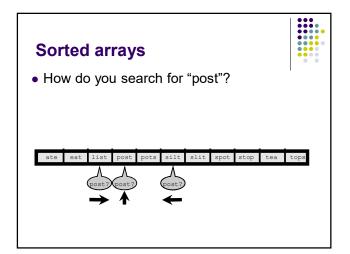
• Insertion is in O(1

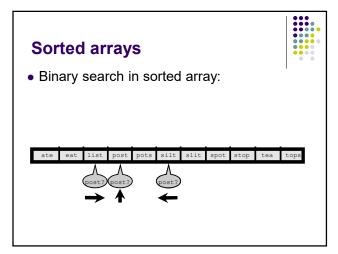


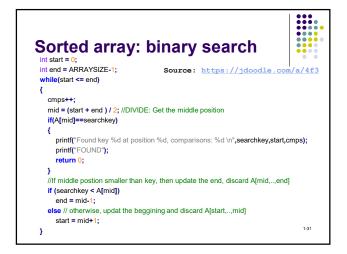




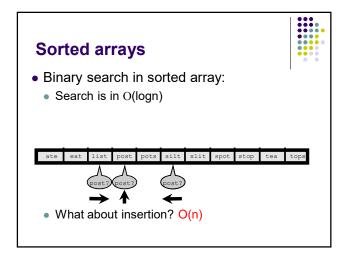
- What is the complexity of search?
  - O(n)

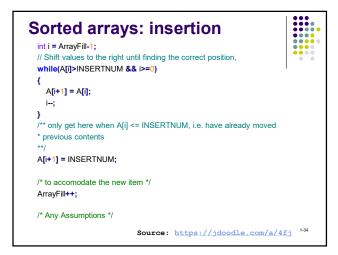


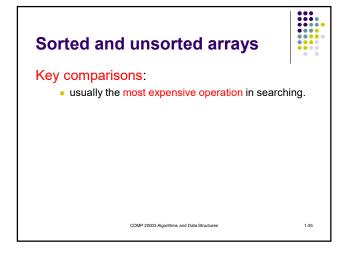


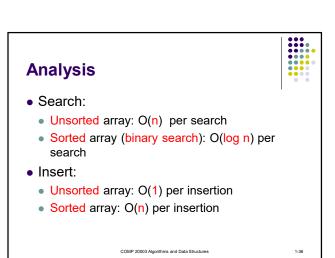


# Searching: analysis • Unsorted array: O(n) • Sorted array (binary search): O(logn)









### Search: m lookups on n items



- Unsorted array, linear search:
- n insertions @ 1 operation → n operations, O(n)
- m lookups @ n operations → m\*n (worst case)
- O(n + m\*n) = O(mn)
- Sorted array, binary search:
  - n insertions @ n comparisons and n data movements each → O(n²)
  - m searches @ log n comps each -> m\*log<sub>2</sub>n
  - O(n² + m log n)

    COMP 20003 Algorithms and Data Structures

### Search: m lookups on n items



- Unsorted array, linear search:
- n + m\*n ~ m\*n
- Sorted array, binary search:
  - n<sup>2</sup> + m log n
- For m << n, unsorted arrays are better!
- But usually m > n, so use sorted array
  - Or even something better!?

COMP 20003 Algorithms and Data Structures

4.00

### Something better?



 What are the worst properties of a sorted array?

COMP 20003 Algorithms and Data Structures

### **Limited size**



- We can overcome the limited size problem using dynamic memory allocation
- C library functions:
  - void \*calloc(size\_t nobj, size\_t size)
  - void \*realloc(void \*p, size\_t size)
  - also, of course void \*malloc(size\_t size)
  - All defined in stdlib.h

OMP 20003 Algorithms and Data Structures

```
malloc(): size_t
malloc(size_t size)
size_t is:

an unsigned integer type
the type returned by the sizeof operator
widely used in the standard library (stdlib) to represent sizes
e.g. malloc(sizeof(int))
```

```
malloc() example (part 1)

#define NUMBER 5
int main (argc, argv)
{
    int var;
    var = NUMBER;
    printf("%d - %d\n", &var, var);
    return 0;
}
>a.out
134509940 - 5
```

# malloc():check return value • Be aware: malloc() can fail! • If malloc() fails, it returns NULL • Never use a pointer to something where the memory allocation has failed! int\* B; B = (int\*) malloc( NUMBER \* sizeof(int)); /\* always check return value of malloc()\*/ If( B == NULL ) { printf("malloc() error\n"); exit(1); } → write a function safemalloc() that does this

```
Getting memory for an array using malloc()

int A[NUMBER];

/**

* while insertions < NUMBER array is OK

* BUT... has a limit

*/

int* B;

/* always check return value of malloc()*/

if( (B = (int*) malloc() NUMBER * sizeof(int) )) == NULL)

{

printf("malloc() error\n");
exit(1);
}

/**

* B can now be used like A

* better to use calloc(NUMBER, sizeof(int))

*/
```

```
Getting memory for an array using calloc()

int* B;

/* always check return value of malloc()*/

if( (B = (int*) calloc( NUMBER, sizeof(int) )) == NULL )
{

printf("malloc() error\n");

exit(1);
}

/* B now comes with each slot initialized to 0 */
```

```
realloc()

/**

* as previously, used malloc(),calloc()

* RESIZE when insertions == NUMBER

*/

B = realloc( B, ( NUMBER * 2 ) * sizeof(int) );

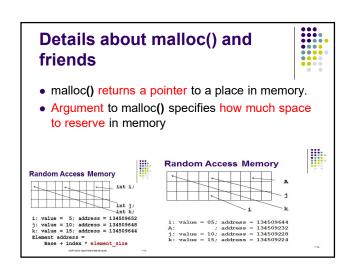
/* should also check realloc() for NULL*/

/* now initialize new part of array */

for(i = NUMBER; i < NUMBER * 2; i++ )

B[i] = NULL;

/* now we have a bigger array, first half copied from the old B */
```



## What's a pointer?



- A pointer is an address in memory
- What is the output of this code?

```
int* ptr;
ptr = (int *)malloc(sizeof(int));
*ptr = 5;
printf("%d, %d", ptr, *ptr);
```

Now what is the output of this code? Try it: https://jdoodle.com/a/4fP

```
ptr = (int*) malloc(sizeof(int));
ptr = 5;
printf("%d, %d", ptr, *ptr);
```

- Int\* p; or int \*p; ? What the inventor of C++ thinks:
  - http://www.stroustrup.com/bs\_faq2.html#whitespace

# Details about malloc() and friends



- #include<stdlib.h>
- Read the documentation for fine points:
  - malloc() returns uninitialized space
  - calloc() returns space initialized to 0
  - realloc(void \*p, size\_t size)
     returns space where the start is copied from p
     and the rest is unintialized.
- Check return value of all memory alloc functions!

COMP 20003 Algorithms and Data Structures

4.50

## malloc() and free()



- malloc() allocates memory.
- free () use to deallocate memory

void\* ptr;
ptr = malloc(NUMBER\_OF\_BYTES);
/\* do things until finished with the
 contents pointed to by ptr \*/
free(ptr);

COMP 20003 Algorithms and Data Structures

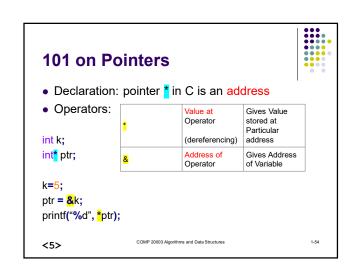
### Back to sorted arrays...



- Space limitations:
- Can use realloc().
- Or can use linked list (sorted linked list).

COMP 20003 Algorithms and Data Structures

# More about Pointers • For an excellent exposition of pointers in C, see the excellent tutorial by Ted Jensen: • LMS Resources → Pointers and Arrays in C



#### 101 on Pointers



- A pointer in C is an address.
- When A is the name of an array, A is a pointer to the array, and
  - A[0] is equivalent to \*(A+0),
  - A[5] is equivalent to \*(A+5)...

20003 Algorithms and Data Structures

## **Memory Allocation: Summary**

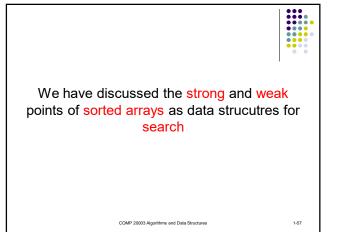


- malloc(),calloc(),and realloc() return:
  - The (untyped) address of allocated memory;
  - i.e. a pointer to allocated memory.

/\*allocates just enough room for an address \*/
struct node\* ptr;

/\* allocates enough room for the node \*/
ptr = (struct node\*) malloc(sizeof(struct node));

COMP 20003 Algorithms and Data Structures



# Linked lists: flexibility, but overheads

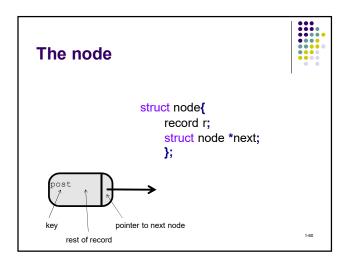


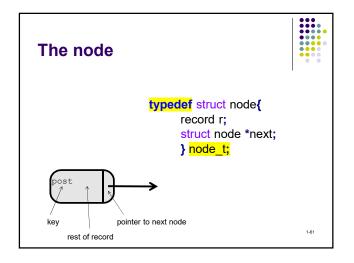
- In a linked list:
  - each item (or key) is located in an arbitrary place in memory,
  - with a link (pointer) to the next item
- Search Operations:
  - If unsorted, finding item is still  $\Theta(n)$ -time.
  - Once insertion point has been determined, easy to insert (or delete) a new item, by rearranging links.

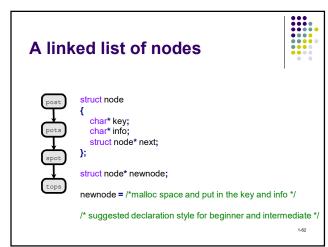
# Linked lists: flexibility, but overheads

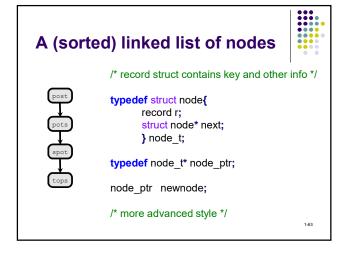


- Takes extra space for each item in the list.
- Takes extra time to allocate the memory for the node for each item.



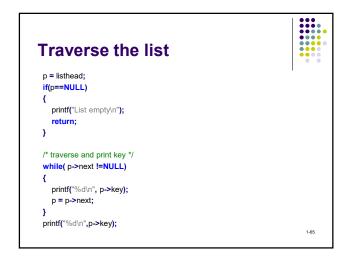


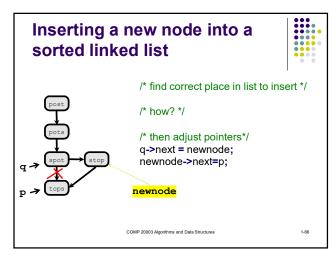


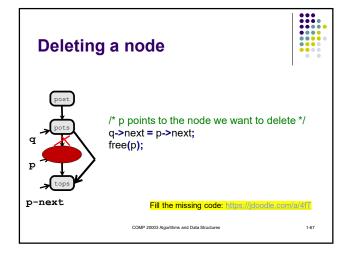


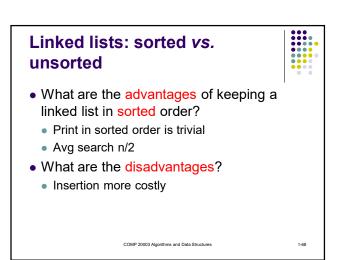
```
Traverse the list

p = listhead;
if(p!=NULL){ /* empty list */
    while( p->next !=NULL)
    {
        printf("%d\n", p->key);
        p = p->next;
    }
    printf("%d\n",p->key);
}
```









### Search: Arrays vs. Linked Lists



- · Sorted arrays:
  - Fast search (binary search), but
  - Slow insertion (keeping sorted order)
- Sorted array:

Sorted linked list

- Fixed size, but
- Can grow with realloc()
- Array needs (in general) only 1 memory allocation
  - Linked list needs many

### Table of "running times"



	One <mark>Search</mark>	One Insert		
Unsorted array	n	1		
Sorted array	log n	n		
Unsorted linked list	n	1		
Sorted linked list	n	n		

#### 

O(n) comps

O(n)

# Practical complexity and algorithms



- O(1): Execute instructions once (or a few times), independent of input
  - Example: pick a lottery winner
- O(log n): keep splitting the input, and only operate on one section of the input.
  - Example: Binary Search
- O(n): Execute instruction(s) once for each item:
  - Example: search in sorted or unsorted array

COMP 20003 Algorithms and Data Structures

# Practical complexity and algorithms



3-73

- O(n log n): split the input repeatedly, and do something to all the segments
  - Example: Many sorting algorithms
- O(n²): For each item, do something to all the others. (Nested loops.)
  - Example:
    - Note: getting slow for large data...
- O(n<sup>3</sup>): matrix multiplication
- O(2<sup>n</sup>): combinatorial problems

COMP 20003 Algorithms and Data Structures

### **Breaking out of linearity**

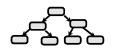


- Compare:
  - Linked list→ → → → → → →
  - Binary tree
- If we reliably know whether the desired item is in the left subtree or the right subtree, we could find it more quickly!

## **Breaking out of linearity**



- Compare:
  - Linked list
  - Binary tree

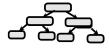


 Note: for a complete binary tree, half the nodes are at the bottom level...

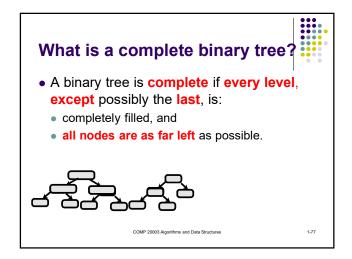
## What is a complete binary tree?

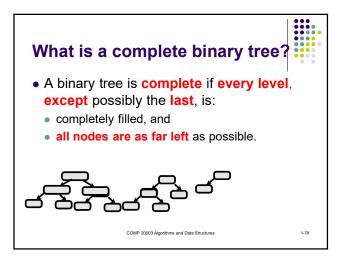


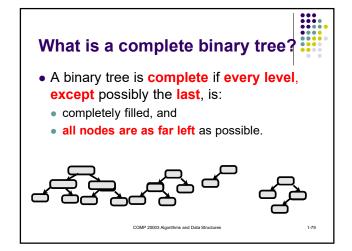
- A binary tree is complete if every level, except possibly the last, is:
  - · completely filled, and
  - all nodes are as far left as possible.

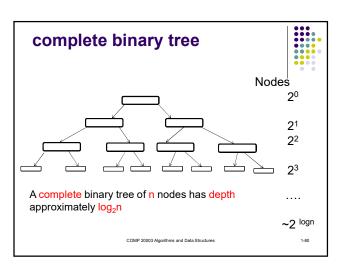


COMP 20003 Algorithms and Data Structures









### **Breaking out of linearity**



- Compare:
  - Linked list
  - Binary tree
- As we will see, both insertion and search are O(log n) operations.

# How a binary search tree work?



- In a sorted linked list:
  - next links to a record with a key ≥ this one.
- In a BST
  - left links to items with key < current key
  - right links to items with key ≥ current key.
- Insert node with key slit in this tree.

# Looking at a complete binary tree Nodes 20 21 22 For a complete binary tree of n nodes, to get to the bottom takes not more than log<sub>2</sub>n key comparisons linked list, to get to the end, we need how many comparisons? n COMP 20003 Algorithms and Data Structures 1-43

### **Binary tree exercizes**



- Put the following numeric keys into a bst:
  - 45, 37, 86, 90, 50, 16, 37
  - How long (how many key comparisons) does it take to search for key=5? 3
- Put the following numeric keys into a bst:
  - 90, 86, 50, 45, 37, 32, 16
  - How long does it take to search for key=5? 7

COMP 20003 Algorithms and Data Structures

# Best case run time in bst: Perfectly balanced tree

- Best case for BST: perfectly balanced
- Height of tree with n items: log<sub>2</sub>n
- Path from root to any node:
  - Maximum length: log<sub>2</sub>n
  - Average length: log<sub>2</sub>n
- Insertion/search/deletion are all O(logn) for a well-balanced tree

# Worst case run time in bst: Stick



- Worst case for BST: a stick.
  - e.g. when items are inserted in sorted order.
  - The BST degenerates to a linked list!
- Height of tree with n items: n
- Path from root to any node:
  - maximum length: n
- average length: n/2
- Insertion/search/deletion are O(n)!

### **Binary search trees**



• Deletion? Next lectures

### Something to think about



 Why don't we just randomize the order of the items we insert into the binary search tree, to prevent worst case behavior?

If you've got the data before do it.

MP 20003 Algorithms and Data Structures

### **Binary search trees**



- Good average case behavior logn
- Bad worst case behavior n
- So overall BST O(n).
  - Actual behavior usually not linear
  - But potentional linearity
- Balanced trees: AVL, red-black; 2,3,4;
   B+tree.

COMP 20003 Algorithms and Data Structures

## **Dictionaries: Summary**



- We have looked at various underlying data structures for implementing dictionaries:
  - •
  - •
  - •

COMP 20003 Algorithms and Data Structures

## **Dictionaries: Summary**



- We have analyzed the computational complexity for these data structures:
  - •

  - •

COMP 20003 Algorithms and Data Structures

### **Dictionaries: Summary**



- So far the best we have done is log n search, where either:
  - Insertion is O(n); or
  - O(log n) average case but O(n) worst case.
- We can do better...

COMP 20003 Algorithms and Data Structures

