


COMP20003
Algorithms and Data Structures
Graph: Shortest Paths

Nir Lipovetzky
 Department of Computing and
 Information Systems
 University of Melbourne
 Semester 2

Example weighted graph

Adjacency List

0→1→2
 1→0→2→3
 2→0→1→3
 3→1→2




Previous visit order from node 0:

- But if these are restaurants and bars, and we want to go to a **nearby** bar From restaurant 0...
- ... in this case the answer is easy. But if you scale it...

COMP 20003 Algorithms and Data Structures

1-2

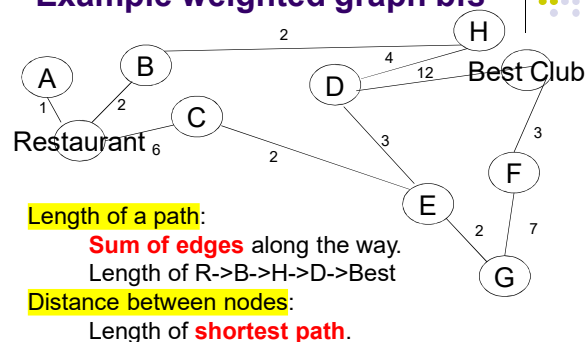
Example weighted graph



COMP 20003 Algorithms and Data Structures

1-3

Example weighted graph bfs



COMP 20003 Algorithms and Data Structures

1-4

Example weighted graph

Questions to ask on a graph:

- Can you **get** from node **X** to node **Y**?
- What is the **shortest distance** from **X** to **Y**?
- What are the **shortest distances** from **X** to **any node**?
- What are the **shortest distances** from **any node** to **any other node**?

1-5

Example weighted graph

Questions to ask on a **directed** graph:

- Can you **get** from node **X** to node **Y**?
- What is the **shortest distance** from **X** to **Y**?
- What are the **shortest distances** from **X** to **any node**?
- What are the **shortest distances** from **any node** to **any other node**?

1-6

Single Source Shortest Path Problem

- Given:
 - Directed graph $G(V, E)$
 - Source vertex s in V
- Determine:
 - Shortest distance** path from s to **every other** vertex in V

COMP 20003 Algorithms and Data Structures 1-7

Brute force approach

- For each vertex v_i :
 - Enumerate all paths from s to v_i
 - Calculate cost of each path $s \rightarrow v_i$
 - Pick minimum cost.
- How many possible paths from s to v_i ?
 - For a dense graph $O(V!)$
 - $V=20$: 2432902008176640000 paths
 - Not feasible!

COMP 20003 Algorithms and Data Structures 1-8

Dijkstra's algorithm for single source shortest path



- **Greedy algorithm:**
 - Based on idea that **any subpath** along a shortest path **is also a shortest path**
 - NodeA $\rightarrow \dots \rightarrow$ NodeX \rightarrow NodeY
 - If **shortest path** from A to Y is **through X**,
 - then this path from A to X is also a **shortest path**
- Dijkstra, E. W., *Numerische Mathematik* 1: 269–271, 1959

COMP 20003 Algorithms and Data Structures

1-9

Dijkstra's algorithm for single source shortest path



- Greedy algorithm:
 - Based on idea that **any subpath** along a shortest path **is also a shortest path**
 - NodeA $\rightarrow \dots \rightarrow$ NodeX \rightarrow NodeY
 - If **shortest path** from A to Y is through X,
 - then this path from A to X is also a **shortest path**
- Assumes **no negative edges**, so:
 - Distance(A \rightarrow X) \leq Distance(A \rightarrow Y)

COMP 20003 Algorithms and Data Structures

1-10

Dijkstra's algorithm for single source shortest path



- Algorithm will give us a **shortest path tree**
- **Root = source node**
 - **Every node** is **connected** to the **root** through its **shortest path**




Image from R. Sedgewick, Lecture Notes
<http://www.cs.princeton.edu/courses/archive/fall05/cos226/lectures/shortest-path.pdf>

COMP 20003 Algorithms and Data Structures

1-11

Dijkstra's Algorithm: Overview



For every vertex v and source s , maintain **estimate** $\text{dist}[v]$ of **minimum distance** $\delta(s, v)$

$\text{dist}[v]$: length of a **known** path $s \rightarrow v$, but not necessarily the shortest path

- $\text{dist}[v] \geq \delta(s, v)$ Always
- When $\text{dist}[v] = \infty$, there is **no estimate** (yet)

Initially $\text{dist}[s] = 0$, all **other** $\text{dist}[v] = \infty$

COMP 20003 Algorithms and Data Structures

1-12

Dijkstra's Algorithm: Overview

Process vertices **one-by-one**, updating **dist[v]** until **dist[v] = $\delta(s,v)$** , for every vertex v

- Along the way, **keep track of best path** information in array **pred[v]**

When algorithm finishes:

- Have shortest distances in **dist[]**
- Can reconstruct shortest path from **pred[]**

COMP 20003 Algorithms and Data Structures

1-13

Relaxation: Updating estimated distances

• Relaxation:

- Estimate** the solution by answering an **easier problem** (relax the conditions)
- dist[]** **Keeps updating** the relaxed estimate **until it is the solution** to the original problem

• For shortest paths:


- Estimate**: known distance of **best path so far**
- Solution**: shortest possible distance

COMP 20003 Algorithms and Data Structures

1-14

Relaxation: Updating estimated distances

• Example:



dist[v]: 35
pred[v]: n


dist[u]: 28
dist[v]: 30
pred[v]: u

COMP 20003 Algorithms and Data Structures

1-15

Relaxation: Updating estimated distances

• Example:




```

void update_relax(u,v)
{
    if( dist[u] + edgeweight(u,v) < dist[v] )
    {
        dist[v] = dist[u] + edgeweight(u,v);
        pred[v] = u;
    }
}
  
```

1-16

Relaxation: Updating estimated distances

- Example:



Note:

```

pred[v] = u;
pred[u] = n;
...
pred[j] = s;

```

Reconstruct path $s \rightarrow v$ going backwards through `pred[]`

1-17

Dijkstra's algorithm: successive relaxations

How do we **pick the next node** to look at?

- Process vertices in order of **estimated closeness** to source, value of `dist[v]`
- **Priority queue** to store vertex v and `dist[v]` value

COMP 20003 Algorithms and Data Structures

1-18

Dijkstra's algorithm (C-ish pseudocode)

```

/* Find shortest paths in graph G from source s */
/* vertices identified by number for convenience */

```

```

void dijkstra(int** G, int s)
{
    int dist[Vsize], pred[Vsize];
    initialize(G, s, pred, dist);
    run(G, s, pred, dist);

    reconstruct(s, pred, dist);
}

```

1-19

Dijkstra's algorithm (C-ish pseudocode)

```

void initialize(int** G, int Vsize, int s, int* pred, int* dist)
{
    int i;
    for( i = 0; i < Vsize; i++)
        dist[i] = MAX_INT;
    dist[s] = 0;
    for( i = 0; i < V; i++)
        pred[i] = NULL;
}

```

COMP 20003 Algorithms and Data Structures

1-20

Dijkstra's algorithm(C-ish pseudocode)

```
void run(int** G, int Vsize, int s, int* pred, int* dist)
{
    pq_node_t* pq;
    int u, v;
    pq = makePQ(G); /* vertices into min PQ, dist as priority */

    while( !emptyPQ(pq) )
    {
        u = deletemin(pq);
        for(/*each v conneted to u */)
            if(dist[u] + edgeweight(u,v) < dist[v])
                update(v, pred, dist, pq);
    }
    /* At this point vertex u has been processed,
    * i.e. dist[u] = delta(s,u) = shortest path to u found */
}
```

1-21

Dijkstra's algorithm
(C-ish pseudocode)

```
void update(int v, int* pred, int* dist, pq_node_t* pq)
{
    dist[v] = dist[u] + edgeweight(u,v);
    pred[v] = u;
    decreaseweight(pq, v, dist[v]);
}
```

1-22

Dijkstra's algorithm
(Example)



D =
Pred =
PQ =

Graph

1-23

Dijkstra's algorithm: example

Example:



Step 0: Initialization.

v	s	a	b	c	d
d[v]	0	∞	∞	∞	∞
pred[v]	nil	nil	nil	nil	nil
color[v]	W	W	W	W	W

Priority Queue:

v	s	a	b	c	d
d[v]	0	∞	∞	∞	∞

Step 1: As Adj[s] = {a,b}, work on a and b and update information.

v	s	a	b	c	d
d[v]	0	2	7	∞	∞
pred[v]	nil	s	s	nil	nil
color[v]	B	W	W	W	W

Priority Queue:

v	a	b	c	d
d[v]	2	7	∞	∞

From lecture notes by Mordechai Golin, Univ Science and Technology, Hong Kong:
<http://www.cse.ust.hk/faculty/golin/COMP271Sp03/Notes/MyL09.pdf>

1-24

Step 2: After Step 1, a has the minimum key in the priority queue. As $Adj[a] = \{b, c, d\}$, work on b, c , and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	10	7
$pred[v]$	nil	s	a	a	a
$color[v]$	B	B	W	W	W

Priority Queue: $\frac{v}{d[v]} \begin{array}{c} b \\ 5 \end{array} \begin{array}{c} c \\ 10 \end{array} \begin{array}{c} d \\ 7 \end{array}$

Step 3: After Step 2, b has the minimum key in the priority queue. As $Adj[b] = \{a, c\}$, work on a, c and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	W	W

Priority Queue: $\frac{v}{d[v]} \begin{array}{c} c \\ 6 \end{array} \begin{array}{c} d \\ 7 \end{array}$

COMP 20003 Algorithms and Data Structures 1-25

Step 4: After Step 3, c has the minimum key in the priority queue. As $Adj[c] = \{d\}$, work on d and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	B	W

Priority Queue: $\frac{v}{d[v]} \begin{array}{c} d \\ 7 \end{array}$

Step 5: After Step 4, d has the minimum key in the priority queue. As $Adj[d] = \{c\}$, work on c and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	B	B

Priority Queue: $Q = \emptyset$.

COMP 20003 Algorithms and Data Structures 1-26

Dijkstra's algorithm (ANALYSIS)

```
void run(int** G, int Vsize, int s, int* pred, int* dist)
{
    pq_node_t* pq; /****** Assuming PQ is a minheap */
    int u, v;
    pq = makePQ(G); /****** big-O() ??? */

    while( !emptyPQ(pq) )
    {
        u = deletemin(pq); /****** big-O() ??? */
        for( /*each v conncted to u */
            if(dist[u] + edgeweight(u,v) < dist[v])
                update(v, pred, dist, pq); /***** big-O() ??? */
    }
}
```

1-27

Dijkstra's Algorithm: Analysis

Cost depends on implementation of PQ

- Using a heap:
 - $makePQ()$ $O(V)$
 - $V * deletemin()$ operations $@O(\log V)$
 - $O(E)$ $decreaseweight()$ ops $@O(\log V)$
 - Total: $O((V+E) \log V)$

1-28

Dijkstra's Algorithm: Limitations

Assumes **no negative edges**:

- Good for physical distances
- Distances are static

Negative edges:

- Use **Bellman-Ford algorithm**
- **Cannot** deal with **negative cycles**
- $O(V \cdot E)$

1-29

Dijkstra's Algorithm: Limitations

Negative cycles:

- What is the **shortest** path?
- Problem is not well-formed, intractable
- Bellman-Ford detects negative cycles (algorithm does terminate, stops keeps shortening paths)

Tutorial:

<https://www.dyclassroom.com/graph/detecting-negative-cycle-using-bellman-ford-algorithm>

1-30

Applications

More applications

- Robot navigation
- Texture mapping
- Typesetting in TeX
- Urban traffic planning
- Optimal pipelining of VLSI chip
- Telemarketer operator scheduling
- Routing of telecommunications messages.
- Network routing protocols (OSPF, BGP, RIP)
- **Exploiting arbitrage opportunities** in currency exchange
- Optimal truck routing through given traffic congestion pattern

1-31

Negative Cycle Detection: Arbitrage

Common example in CS materials is arbitrage:

- currency 1 \rightarrow currency 2 \rightarrow currency 3 \rightarrow currency 1'
- If currency 1' > currency 1, you have made money
- Model problem as a graph:
 - Vertices = currency
 - Edges = $-\log_2(\text{exchange rate})$
 - Detect negative cycle and change money \rightarrow get rich!

Not realistic!

- D.J.Fenn *et al.*, "The Mirage of Triangular Arbitrage in the Foreign Currency Exchange Market", *Int. J. Theoretical and Applied Finance* 12(8), 1105-1123, 2009.

1-32

Edsger W. Dijkstra

- The question of whether **computers can think** is like the question of whether **submarines can swim**
- **Computer science** is no more about **computers** than **astronomy** is about **telescopes**
- How do we convince people that in programming **simplicity** and **clarity** —in short: what mathematicians call "**elegance**"— are not a dispensable luxury, but a crucial matter that decides between success and failure?

Elegance is not a dispensable luxury but a quality that decides between success and failure

Turing award 1972

1-33

A very nice explanation of Dijkstra's algorithm by Mordechai Golin can be found at

<http://www.cse.ust.hk/faculty/golin/COMP271Sp03/Notes/MyL09.pdf>

COMP 20003 Algorithms and Data Structures

1-34